


# Binary Search Trees


# Definitions

- A *tree* is a collection of *nodes* and directed edges such that
  - Each node has a unique parent node, except the root node, which has no parent
  - Each node has a (possibly empty) set of children
  - Parent and child nodes are connected by a directed edges from parent to child
  - There is a unique path from the root to each node in the tree



# Definitions

- A *leaf* is a node that has no children
- An *internal node* is a node that has at least one child
- Two nodes with the same parent are called siblings
- *Grandparent, ancestor, descendant* etc. are defined analogously
- A *subtree* consists of a node along with *all* its descendants




# Definitions

- A *path* from node  $n_0$  to node  $n_k$  is a sequence  $n_0, n_1, \dots, n_k$  such that for all  $i$ ,  $n_i$  is the parent of  $n_{i+1}$
- The *length* of the path is the number of *edges* it contains
  - Not the number of nodes (usually)
- For each node  $n$  the *depth* of  $n$  is the length of the unique path from the root to  $n$
- The *height* of  $n$  is the length of the longest path from  $n$  to a leaf
- The *height* of a tree is equal to the height of the root node of the tree

# Definitions

- The *branching factor* of a tree is the maximum number of children in any of its nodes
- A *binary tree* is a tree with a branching factor 2
- A *full node* has the maximum number of children
- A *level* in a tree is the set of all nodes in the tree at a given depth
- A *complete tree* is one where all levels are full except the bottom level, which has been filled from left to right
  - Heaps are complete trees
- A *full tree* is a complete tree whose last level has been filled completely

# Example



- The path c-f-k has length 2
- The (depth of f) = (depth of g) = 2
- (height of f) = 1, (height of g) = 0
- (depth of a) = 0
- The height of the tree is 3

# Traversal

- Want to visit all the nodes in the tree in some order, and apply some operation to each node

## preorder traversal pseudocode


```
void pre_order_traversal(node)
{
    visit(node);
    for each child of node {
        pre_order_traversal(child);
    }
}
```

## postorder traversal pseudocode

```
void post_order_traversal(node)
{
    for each child of node {
        post_order_traversal(child);
    }
    visit(node);
}
```

## preorder traversal pseudocode

```
void pre_order_traversal(node)
{
    visit(node);
    for each child of node {
        pre_order_traversal(child);
    }
}
```




Pre-order traversal (aka depth-first traversal): a, b, e, c, f, j, k, g, d, h, i

- Idea: visit node, then visit each of its subtrees. When visiting a subtree, start from the root



## postorder traversal pseudocode

```
void post_order_traversal(node)
{
    for each child of node {
        post_order_traversal(child);
    }
    visit(node);
}
```




Post-order: e, b, j, k, f, g, c, h, i, d, a

Visit each of the subtrees of the node, then visit the node itself

- For binary trees (each node has at most two children): *in-order traversal*

```
void in_order_traversal(node)
{
    if (node->left)    in_order_traversal(node->left);
    visit(node);
    if (node->right)   in_order_traversal(node->right);
}
```




In-order traversal: h, d, i, b, e, a, f, c, g

# Visualizing the call stack

```
int x = f3(f2(f1()), f3(f1()), f2(5)+1)
        6   2   1   5   3       4
```

We can construct a tree where f2 is the child of f3 is the call to f2 is needed to evaluate f3, etc.

Traverse the tree in postorder to determine the order of calls



# Tree implementation

```
typedef struct bin_tree_node {  
    const void *data;  
    struct bin_tree_node *left;  
    struct bin_tree_node *right;  
} bin_tree_node_t;  
  
typedef struct bin_tree {  
    bin_tree_node_t *root;  
    size_t size;           /* not strictly necessary */  
} bin_tree_t;
```

# Searching a tree

- Want to find a node with a specific value in the tree
- $n$  nodes
- Worst-case run-time:  $O(n)$

```
search(root, value)
    if root = NULL
        return NULL
    if root->data = value
        return root
    left = search(root->left)
    if(left != NULL)
        return left
    return search(root->right)
```

```
bool search(const bin_tree_node_t *root, const void *value)
{
    return root != NULL &&
        (root->data == value ||      /* think about this... */
         search(root->left, value) ||
         search(root->right, value));
}
```

# Binary Search Trees (BST)


- A tree  $T$  where
  - All values in  $T$  are *comparable* (can be ordered)
  - All nodes in the left subtree of a node with value  $k$  have a value less than  $k$
  - All nodes in the right subtree of a node with value  $k$  have a value greater than  $k$
  - (Implication: all values are different)

# Searching a BST

```
search(root, value)
    if root = NULL
        return NULL
    if value = root->data
        return root
    else if value < root->data
        return search(root->left, value)
    else
        return search(root->right, value)
```

# BST search run-time

- In the worst case,  $O(n)$
- Need to traverse all nodes
- If height  $h$  is known, runtime is  $O(h)$
- If the tree is complete, the height is approx.  $\log(n)$ , so the runtime is  $O(\log(n))$





# BST implementation

- Use the same structure as for an ordinary binary tree, but add a comparator function


```
typedef struct bst_node {
    const void *data;
    struct bst_node *left;
    struct bst_node *right;
} bst_node_t;

struct bst {
    bst_node_t *root;
    size_t size;
    int (*cmp)(const void *, const void *);
};
```


# Inserting into a BST

- Need to insert a new node into the BST, and want for the resulting tree to still be a BST

```
insert(cur_node, value)
    if cur_node = NULL
        create a node with data=value, insert it instead of cur_node
    if value < cur_node->data
        insert(cur_node->left, value)
    else
        insert(cur_node->right, value)
```




insert 4  
→




# Insertion order matters

- Insert 3, 1, 0, 2, 5, 4, 6



# Insertion order matters


- Insert 1, 2, 3, 4, 5, 6



# Deleting from a BST


- Plan:
  - Find the node to delete
  - Delete it
- Cases of node deletion
  - Delete a node with no children (easy)
  - Delete a node with one child
  - Delete a node with two children

# Deleting a node with no children



# Deleting a node with one child

- Replace the node with its child
- The child has the same properties as the parent with respect to the grandparent – if the parent is smaller than the grandparent, so is the child; if the parent is larger than the grandparent, so is the child





# Deleting a node with two children

- Leave the node in the tree
- Replace the value stored in the node with the successor value
  - Successor: the next larger value in the tree
- Idea: the successor is larger or equal to the left child, and smaller or equal to the right child

# Finding the successor

```
successor(node)
```

```
    cur_node = node->right
```

```
    while cur_node->left != NULL
```


```
        cur_node = cur_node->left
```

```
    return cur_node
```

- Idea: look in the right subtree, and then go left as much as possible
  - Need the smallest node in the right subtree

# Deleting a node with two children

- Find the successor
- Swap the node with its successor
- Delete the node from its new location
  - The successor has at most one child, since it definitely doesn't have a left subtree
    - We went left as much as possible



# BST Implementation

- Bag ADT
  - A collection of elements
  - Want to make the ADT independent of the implementation

```
bag_t *bag_create(int (*cmp)(elem_t, elem_t));
```

```
void bag_destroy(bag_t *b);
```

```
size_t bag_size(const bag_t *b);
```

```
bool bag_contains(const bag_t *b, elem_t e);
```

```
bool bag_insert(bag_t *b, elem_t e);
```

```
bool bag_remove(bag_t *b, elem_t e);
```

# BST implementation of Bag

In bst\_bag.c and not bag.h since those structs are only used internally

```
/* TYPE bst_node_t -- The type of tree nodes used to store elements. */
typedef struct bst_node {
    elem_t elem;
    struct bst_node *left;
    struct bst_node *right;
} bst_node_t;

/* TYPE bag -- Definition of struct bag from the header file. */
struct bag {
    size_t size;
    bst_node_t *root;
    int (*cmp)(elem_t, elem_t);
};
```

# Bag functions call BST functions

(not declared in bag.h)

```
void bst_destroy(bst_node_t *root)
{
    if (root) {
        // Standard post-order traversal to free all the memory.
        bst_destroy(root->left);
        bst_destroy(root->right);
        free(root);
    }
}
```

(declared in bag.h)

```
void bag_destroy(bag_t *b)
{
    // Recursively free the memory for tree nodes, then free b itself.
    bst_destroy(b->root);
    free(b);
}
```

# Bag functions call BST functions

```
bool bst_remove(bst_node_t **root, elem_t elem, int (*cmp)(elem_t, elem_t))
{
    // This function takes a *pointer* to the pointer to the root node, so that
    // we can modify its value directly -- for example, when this is called with
    // argument &b->root, it gets the *address* of the pointer b->root, so that
    // if b->root->elem == elem and b->root is the only node, we can change the
    // value of b->root to become NULL directly from here.
    if (! *root)
        return false;
    else if (cmp(elem, (*root)->elem) < 0)
        return bst_remove(&(*root)->left, elem, cmp);
    // ...
}

bool bag_remove(bag_t *b, elem_t e)
{
    // Remove the element recursively, starting at the root.
    if (bst_remove(&b->root, e, b->cmp)) {
        b->size--;
        return true;
    } else {
        return false;
    }
}
```

# remove\_min

- If root doesn't have a left child, it is the minimum
  - Replace the root with its right child
- If root has a left child, remove\_min(left child)

```
elem_t bst_remove_min(bst_node_t **root)
{
    // As above, this function takes a pointer to a pointer to the root node, so
    // that the value of the pointer can be modified directly from in here.
    if ((*root)->left) {
        /* *root is not the minimum, keep going */
        return bst_remove_min(&(*root)->left);
    } else {
        /* remove *root */
        bst_node_t *old = *root;
        elem_t min = (*root)->elem;
        *root = (*root)->right;
        free(old);
        return min;
    }
}
```



# bst\_remove

```
bool bst_remove(bst_node_t **root, elem_t elem, int (*cmp)(elem_t, elem_t))
{
    // This function takes a *pointer* to the pointer to the root node, so that
    // we can modify its value directly -- for example, when this is called with
    // argument &b->root, it gets the *address* of the pointer b->root, so that
    // if b->root->elem == elem and b->root is the only node, we can change the
    // value of b->root to become NULL directly from here.
    if (! *root)
        return false;
    else if (cmp(elem, (*root)->elem) < 0)
        return bst_remove(&(*root)->left, elem, cmp);
    else if (cmp(elem, (*root)->elem) > 0)
        return bst_remove(&(*root)->right, elem, cmp);
    else { /* (cmp(elem, (*root)->elem) == 0) */
        // We've found the value to remove; check if *root has two children.
        if ((*root)->left && (*root)->right) {
            (*root)->elem = bst_remove_min(&(*root)->right);
        } else { /* remove *root */
            bst_node_t *old = *root;
            *root = (*root)->left ? (*root)->left : (*root)->right;
            free(old);
        }
        return true;
    }
}
```

# Using Bag

```
int float_cmp(const void *a, const void *b)
{
    return *(float *) a < *(float *) b ? -1
        : *(float *) a > *(float *) b; /* ? 1 : 0 redundant */
}
```

```
bag_t *b1 = bag_create(float_cmp);
```

```
bag_insert(b1, &elts[i])
```